

Back-propagation Tutorial

Mingwen Dong

December 28, 2016

1 Backpropagation

Notations:

1. z_i^l : weighted linear (sum) input to i^{th} neuron in layer (l) .
2. a_i^l : activation/output from i^{th} neuron in layer (l) .

$$a_i^l = g(z_i^l)$$

$g(\cdot)$ is the activation function. e.g., sigmoid function $g(x) = \frac{1}{1+e^{-x}}$, $\tanh(\cdot)$, or ReLU.
Note: activation function could be other form like Huber's function.

3. w_{ji}^l : connection strength/weight from i^{th} neuron in layer $(l-1)$ to j^{th} neuron in layer (l) .

$$z_j^{l+1} = \sum_{i=0}^n w_{ji}^{l+1} a_i^l$$

$$\frac{\partial z_j^{l+1}}{\partial z_i^l} = \frac{\partial z_j^{l+1}}{\partial a_i^l} \frac{\partial a_i^l}{\partial z_i^l} = w_{ji}^{l+1} g'(z_i^l)$$

Notice: the bias is included by adding a constant input "1" to every neuron in the network.

4. L : the output layer.
5. C : the cost. e.g., quadratic cost function, cross-entropy, or likelihood cost function.

$$\frac{1}{2} \sum_{i=0}^n (a_i^L - y_i)^2$$

$i = 0, 1, 2, \dots, n$, indicate different output units/neurons.

In stochastic gradient descent, the cost is also summed over different input x .

6. δ_i^l : derivative of the cost C with respect to each neuron's linear input z_i^l

$$\delta_i^l = \frac{\partial C}{\partial z_i^l}$$

1.1 Feed Forward

The weighted input to neurons in layer (l) is:

$$z_j^l = \sum_{i=0}^n a_i^{l-1} w_{ji}^l \Rightarrow \tilde{z}^l = \mathbf{W}^l \cdot \tilde{a}^{l-1}$$

where, \mathbf{W}^l is the weight matrix for neurons in layer (l). "tilde" indicate a vector variable. The corresponding activations from these neurons are:

$$\tilde{a}^l = g(\tilde{z}^l)$$

The derivative of a_i^l with respect to z_i^l can be calculated during this forward pass process:

$$\frac{\partial a_i^l}{\partial z_i^l} = g'(z_i^l) = \begin{cases} g(z_i^l)[1 - g(z_i^l)] & \text{sigmoid activation function} \\ 1 & \text{if } z_i^l > 0 \quad \text{otherwise } 0 \end{cases}$$

1.2 Errors at output layer

Define cost-function as cross-entropy:

$$C = \sum_{i=0}^n \left[y_i \ln(a_i^L) + (1 - y_i) \ln(1 - a_i^L) \right] \quad i = 0, 1, \dots, n \text{ indicate different neurons in the output layer}$$

Errors at the output layer (one could treat cost as another neuron where all output neurons converge to):

$$\delta_i^L \equiv \frac{\partial C}{\partial z_i^L} = \frac{\partial C}{\partial a_i^L} \frac{\partial a_i^L}{\partial z_i^L} = \left[\frac{y_i}{a_i^L} - \frac{1 - y_i}{1 - a_i^L} \right] \odot g'(z_i^L) \quad \odot: \text{pointwise multiplication}$$

1.3 Back-propagate errors from output layer to hidden layers

Proof for the back-propagation algorithm using Dynamic Programming.

Using multivariate chain rule, we have:

$$\begin{aligned} \delta_c^l &= \frac{\partial C}{\partial z_c^l} \quad \text{sum over all possible product sequences from layer } (l+1) \text{ to output layer } (L) \\ &= \sum_{i,j,k,m,\dots} \frac{\partial C}{\partial z_m^L} \frac{\partial z_m^L}{\partial z_k^{L-1}} \cdots \frac{\partial z_k^{l+3}}{\partial z_j^{l+2}} \frac{\partial z_j^{l+2}}{\partial z_i^{l+1}} \frac{\partial z_i^{l+1}}{\partial z_c^l} \\ &= \sum_i \frac{\partial z_i^{l+1}}{\partial z_c^l} \cdot \sum_{j,k,m,\dots} \frac{\partial C}{\partial z_m^L} \frac{\partial z_m^L}{\partial z_k^{L-1}} \cdots \frac{\partial z_k^{l+3}}{\partial z_j^{l+2}} \frac{\partial z_j^{l+2}}{\partial z_i^{l+1}} \\ &= \sum_i \frac{\partial z_i^{l+1}}{\partial z_c^l} \cdot \delta_i^{l+1} \\ &= \frac{\partial a_c^l}{\partial z_c^l} \sum_i \frac{\partial z_i^{l+1}}{\partial a_c^l} \cdot \delta_i^{l+1} \\ &= g'(z_c^l) \sum_i w_{ic}^{l+1} \cdot \delta_i^{l+1} \\ &= g'(z_c^l) \left\{ [\tilde{w}_{\cdot c}^{l+1}]^T \cdot \delta_i^{l+1} \right\} \quad \tilde{w}_{\cdot c}^{l+1}: c^{th} \text{ column of weight matrix } \mathbf{W}^{l+1} \end{aligned}$$

This leads to an recursion and in the output layer, we have:

$$\delta_i^L \equiv \frac{\partial C}{\partial z_i^L} = \left[\frac{y_i}{a_i^L} - \frac{1 - y_i}{1 - a_i^L} \right]$$

Write in vectorized computation:

$$\tilde{\delta}^l = g'(\tilde{z}^l) \odot \left\{ [\mathbf{W}^{l+1}]^T \cdot \tilde{\delta}^{l+1} \right\}$$

If we rewrite the recursion as a loop starting from the base case, it's the **back-propagation algorithm**. Intuitively, we could think the errors at the output layer as a new "input" (δ^L), the errors at the hidden layer are obtained by backwards multiplying (δ^L) with transposed weight matrix $[\tilde{w}^{l+1}]^T$ and then scaled with the neuron-specific derivatives. To some extent, this calculation is even simpler either than the feed-forward pass as the calculation only involves elementwise multiplication (assuming derivative is already calculated in the feed-forward process).

1.4 Weight update

From the computation above, we know:

$$\begin{aligned}\frac{\partial C}{\partial w_{ji}^l} &= \frac{\partial C}{\partial z_j^l} \cdot \frac{\partial z_j^l}{\partial w_{ji}^l} \\ &= \delta_j^l \cdot a_i^{l-1}\end{aligned}$$

Write in vectorized form:

$$\begin{aligned}\frac{\partial C}{\partial \tilde{w}_j^l} &= \delta_j^l \cdot \tilde{a}^{l-1} \\ \Rightarrow \frac{\partial C}{\partial \mathbf{W}^l} &= \tilde{\delta}^l \cdot [\tilde{a}^{l-1}]^T\end{aligned}$$

If learning rate is η , the new \mathbf{W}^l should be:

$$\mathbf{W}^l \leftarrow \mathbf{W}^l - \eta \cdot \frac{\partial C}{\partial \mathbf{W}^l} = \mathbf{W}^l - \eta \cdot \tilde{\delta}^l \cdot [\tilde{a}^{l-1}]^T$$

1.5 Stochastic gradient descent

Forward pass:

1. \mathbf{X} : input matrix, dimension is $m \times n$ (n examples each with m features), each column indicates one example.
2. \mathbf{W} : weight matrix, dimension is $p \times m$ (receives from m inputs, output p linear sum).
3. $\mathbf{Z} = \mathbf{W} \cdot \mathbf{X}$: output matrix, dimension is $p \times n$.

Backward pass:

1. δ^L : error matrix at the output layer, dimension is $p \times n$, each column indicates output error from one example.
2. \mathbf{W} : weight matrix, dimension is $p \times m$, each row indicates the connection from one input neuron to p output neuron.
3. δ^l : error matrix at the hidden layer, dimension is $m \times n$, each column indicates the errors at hidden layer (l) from one example.

2 Regularization

For L2 regularization (assume the regularization strength is λ), the cost function (cross-entropy) is:

$$C = \sum_{i=0}^p \left[y_i \ln(a_i^L) + (1 - y_i) \ln(1 - a_i^L) \right] + \frac{\lambda}{2} \sum_{i,j,l} [w_{ij}^l]^2 \quad \text{all weights in the network}$$

$$\frac{\partial}{\partial w_{ij}^l} \left\{ \frac{1}{2} \sum_{i,j,l} [w_{ij}^l]^2 \right\} = w_{ij}^l$$

The derivative from regularization term is directly related to each weight and doesn't need back-propagation. For stochastic gradient descent, the update rule becomes (the regularization term is scaled again by the size of training set n):

$$\mathbf{W}^l \leftarrow \mathbf{W}^l - \frac{\eta\lambda}{n}\mathbf{W}^l - \eta \cdot \frac{\partial C}{\partial \mathbf{W}^l} = \mathbf{W}^l - \frac{\eta\lambda}{n}\mathbf{W}^l - \eta \cdot \delta^l \cdot [\tilde{a}^{l-1}]^T$$